

Effects of Cooperative and Competitive Coevolution on Complexity in a Linguistic Prediction Game

Nick Moran¹ and Jordan Pollack¹

¹DEMO Lab, Brandeis University, 415 South St, Waltham, MA
nemti@brandeis.edu

Abstract

We propose a linguistic prediction game with competitive and cooperative variants, and a model of game players based on finite state automata. We present a complexity metric for these automata, and study the coevolutionary dynamics of complexity growth in a variety of multi-species simulations. We present quantitative results using this complexity metric and analyze the causes of varying rates of complexity growth across different types of interactions. We find that while both purely competitive and purely cooperative coevolution are able to drive complexity growth above the rate of genetic drift, mixed systems with both competitive and cooperative interactions achieve significantly higher evolved complexity.

Introduction

Explaining the causes and mechanisms of the evolution of increasing complexity in biological systems, and reproducing those mechanisms *in silico* remains an open problem for the field of artificial life (Bedau et al., 2000; Taylor et al., 2016). A variety of hypotheses have been proposed in both the biological literature at large, and within the evolutionary computation community. In his book *Full House* (Gould, 2011), biologist Stephen J. Gould suggests that the apparent increase in biological complexity over evolutionary time is explained by the process of random mutation performing a drunkard's walk through the space of possible phenotypes. Because there is a lower bound of complexity at zero, this random walk will, over time, reach higher and higher maxima. On the other hand, Richard Dawkins has offered co-evolutionary "arms races" to explain the development of ever-higher complexity (Dawkins and Krebs, 1979). In this hypothesis, co-evolutionary competition drives a positive feedback cycle of complexity growth, in which each new innovation by one competitor drives selection for further innovation in the other.

There are a variety of challenges when examining this sort of question, but perhaps chief among them is the challenge of satisfactorily defining complexity. The literature contains a wide variety of definitions in use across computer science (Wolfram et al., 1984; Adami, 1998; Dorin and Korb, 2010), biology (Adami et al., 2000; Dawkins, 2010), and in

other fields. A common thread in many of these definitions (though certainly not all), is the notion of incompressibility - that is, that things are complex when they cannot be succinctly described. In particular, the concept of algorithmic complexity developed by Kolmogorov and Chaitin defines the complexity of a string as the length of the shortest program on a universal Turing machine which generates that string (Kolmogorov, 1968; Rissanen, 1978). Unfortunately, this measure is difficult to use in practice, as it cannot be computed in the general case. In the field of automata theory, the state complexity of a regular language is defined as the size of the smallest finite automaton which can recognize that language (Yu, 2001). This measure is conceptually similar to that of Kolmogorov and Chaitin, but can be calculated by a straightforward algorithm. We will present a variant of this metric, and use it to evaluate the complexity of our evolved phenotypes.

In this work, we propose a simple model to serve as a testbed for these and other hypotheses. This model consists of a co-evolutionary game between multiple species, a representation of strategies for this game, and a complexity measure over that representation. The game is a two-player, iterated game in which each individual generates a binary string, and receive fitness points depending on whether their string matches or differs from their opponent's string at each index. This game has both a competitive form, in which the two players have opposing goals (that is, one seeks to match and the other seeks to differ), and a cooperative form, in which the two players share a goal (that is, both seek to match, or both seek to differ). We model players of this game as finite state machines, which have been extended to label each state with an output symbol. We then calculate the state complexity of a strategy by applying a minimization algorithm to find the smallest machine which can produce that strategy. Further details on each of these components can be found in their corresponding sections.

Using this model, we can construct a variety of ecological networks, in which we define either competitive or cooperative pair-wise relationships between multiple species. In networks with more than two species, a single species may

have multiple such relationships, some competitive and others cooperative, and will derive fitness from all of them. We then compare the growth of complexity of various network topologies in order to determine which, if any, promote complexity, and to which degree.

Model

Linguistic Prediction Game

We present a simple prediction game which we term the Linguistic Prediction Game. In this game, two players iteratively generate symbols (in this work, 0 or 1) and are rewarded points each round depending on their symbol and the symbol of the other player. Symbol generation in each round is done simultaneously - neither player may know what symbol the other has chosen to produce. However, players may condition their next symbol on the history of symbols played so far by both players in the current game.

Many different payoff matrices can be defined, but we concern ourselves with a few simple matrices which model cooperation and competition. In these games, each player either seeks to produce a symbol which matches that of the other (in other words, both produce zero or both produce one), or seeks to produce a symbol which mismatches that of their opponent (in other words, one player produces zero and the other produces one)¹. If the goals of both players are aligned, meaning they both wish to match or both wish to mismatch, then we say the game is cooperative, and if their goals are opposed, then we say the game is competitive.

In each round, a player receives one point if their goal (match or mismatch) succeeds, and zero points otherwise. In this work, we consider a game which is played over an infinite number of rounds (similar to the formulation of an infinite iterated prisoner's dilemma in (Lindgren, 1992)), and we define the score of a particular player via a limit:

$$S(p) = \lim_{r \rightarrow \infty} \frac{\sum_{i=0}^r \text{reward}(p, i)}{r}$$

Where $S(p)$ indicates the total score of player p , r the number of rounds played, and $\text{reward}(p, i)$ indicates the reward received by player p in round i .

We consider two specific types of game matrices, as shown in Figure 1, which we term COOP and COMP. COOP is a cooperative game in which both players receive a point if they output the same symbol. In this game, players must learn to coordinate their outputs to receive a maximal score. Players seek to establish a shared convention (a language) to which they can both adhere. COMP, on the other hand, is a competitive game in which one player receives a point if they output the same symbol, and the other receives a point if they output different symbols. Here both players must try

¹These games may be viewed as variants of the iterated penny-matching game described in (Gibbons, 1992) and further explored in (Mookherjee and Sopher, 1994)

	0	1
0	1,1	0,0
1	0,0	1,1

COOP

	0	1
0	1,0	0,1
1	0,1	1,0

COMP

Figure 1: Cooperative and competitive game matrices. Symbols output by the first player are on the left axis, and those output by the second player along the top axis. Each cell is labeled with the payoffs received by the first player and the second player.

to predict the pattern of their opponent while simultaneously preventing prediction of their own pattern.

Note that there are two additional isomorphic formulations of each game - COOP also occurs when both players seek to output different symbols, and COMP also occurs when the roles of the two players are swapped. In our experiments, the choice of which form of each game to use has had no effect on the results.

Multi-game Ecosystems

We extend this linguistic prediction game model by embedding multiple instances of it in a network of interacting species, which we term an ecosystem. An ecosystem consists of a set of species, and a set of pairwise connections between them. Each connection is labelled with a game matrix which governs the interaction between members of those two species. In this way, a single organism may interact with other organisms of multiple different species, and will potentially play a different version of the game with each of those species.

We examine several of these ecosystem models here, as shown in Figure 2. The first is a trivial baseline model with a single species and no games being played. This serves as a control datapoint which allows us to understand the complexity which tends to arise only from a random walk through mutations. The next two models are a pair of two-species models which are fully cooperative (playing the COOP game), and fully competitive (playing the COMP game). These two models allow us to examine the behavior that occurs when the players have either aligned or opposing goals. We also consider an ecosystem which features a mix of cooperative and competitive games, which we term "Three-Species Mixed". This ecosystem consists of three species, A , B and C . A plays the COOP game against B




Figure 2: Interactions between species in our six ecosystem models. Connections between species indicate an interaction, which is labeled as either competitive or cooperative. For competitive interactions, the arrow points from the species which is rewarded for matching to the species which is rewarded for mismatching. Bidirectional arrows indicate cooperative interactions.

and the COMP game against C , and there is no direct interaction between B and C . This system is inspired by the three-player pursuit/evasion game introduced by Ficici (Ficici et al., 1998). For comparison with this three-species system, we also examine a purely competitive three-species system in which A competes with B and with C , but with opposite goals - that is, A attempts to match B but mismatch C . We term this system “Three-Species Comp”. Finally, we consider a four-species system extension of Ficici’s game in which both species A and B compete, and each has a cooperative partner, C and D respectively. A and C play COOP as do B and D . We term this system “Four-Species Mixed”.

Critically, in ecosystems in which one individual may interact with multiple other species, there is no signal or indication of which such other species is being interacted with. For example, in the Three-Species Mixed system, individuals of species A will not know ahead of time whether they are interacting with a member of species B , and therefore playing COMP, or a member of species C , and therefore playing COOP. To the extent that a member of species A exhibits different behavior in these two situations, it must

do so by learning to recognize the species of its opponent based on its opponent’s outputs, and adjust accordingly.

Organism Model

Individual organisms are represented as finite state machines in which each state is labelled with a symbol to be produced (either zero or one), and each state has an outgoing transition link for each symbol. State transitions occur according to the symbols produced by the other player in game. Additionally, each finite state machine has a special initial state in which it will start at the beginning of each game. Thus, an organism can be represented as a tuple S, T, L, s_0 , where S is a set of states, $T : S \times \{0, 1\} \mapsto S$ defines state transitions for each state in S for both 0 and 1, $L : S \mapsto \{0, 1\}$ defines the output symbol for each state in S , and s_0 indicates the initial state.

To simulate a game, each player is first set to its initial state $s_0 \in S$. Then, each player produces the symbol indicated by the label on its current state, and performs the state transition indicated by the symbol produced by the opponent. To calculate the score received by each player, this process is repeated until the two players enter a loop (note that players are fully deterministic, and possess a finite number of internal states, and therefore a pair of players must enter a finite-length loop after a finite number of rounds). To detect loops, we track the pairs of states of each player in each round, and wait for the system to return to the same pair of states as it has been in before. Each player’s score for the game is then its average score over the loop. Note that this means interactions before the loop is entered (i.e. the transient state of the system prior to its stable cycle) do not contribute to the score, as they contribute an infinitesimal amount in the limit of an infinite game.

As an example, consider the two organisms pictured in Figure 3. Let the first organism be A , with states A_0 and A_1 , and the second be B with states B_0 and B_1 . A game between these two players will begin in states $\{A_0, B_0\}$, and will proceed as $\{A_0, B_0\} \rightarrow \{A_1, B_0\} \rightarrow \{A_1, B_1\} \rightarrow \{A_0, B_1\} \rightarrow \{A_1, B_0\}$. At this point, the game has entered a loop, and the score can be calculated using a game matrix.

Mutation We further define a set of mutation operators which alter the structure of an organism². These are:

- Add a new state.
- Remove a state.
- Redirect a transition link.
- Change the symbol produced in a state.

²These mutation operators are similar to the class of structural mutators used in (Angeline et al., 1994) and other network-based evolutionary algorithms.




Figure 3: Two examples of organisms. The double circle represents the initial state, the node labels indicate the symbol to be produced when in that state, and the transition labels indicate which symbol causes that transition to occur. The first organism will output 0 on the first round, and then copy the opponent's last move. The second organism will also output 0 on the first round, and then produce the opposite of the opponent's last move.




Figure 4: All-vs.-all interactions between two species. Each member of species A interacts with every member of species B, but two members of the same species do not interact with one another.

Adding a new state results in a state with randomly-assigned outgoing transitions (possibly pointing to itself), and with a random output symbol. Note that this new state will not yet be the target of any external transitions - it must later be linked by redirection mutation. Removing a state requires randomly selecting new destinations for every link which previously pointed to the deleted node. Further, if the deleted node was the initial state, a new initial state is randomly selected from the remaining nodes. If the organism has only one state, a deletion mutation instead has no effect. A redirection mutation simply selects a random transition link and selects for it a new destination. This may result in making some nodes unreachable, or making some previously unreachable nodes reachable. Changing the symbol produced by a state is self-explanatory.

Evolutionary Simulation

An evolutionary simulation consists of the following steps. First, for each species in an ecosystem, a population of organisms is initialized with minimal finite state machines - a single state with a random label, and transition links leading back to that single state. We then simulate a series of generations. Each generation consists of an evaluation period, in which individuals interact according to the games specified by the ecosystem to accumulate fitness, and a reproduction period in which individuals produce offspring which form

the populations of the next generation.

In the evaluation period, for each pair of interacting species, an all-vs.-all set of games are played, as shown in Figure 4. That is, if species *A* plays the COOP game with species *B*, each member of *A* will play that game with each member of *B*. For each organism, its total fitness across all interactions is tallied.

In the reproduction stage, for each species, a new population is generated using fitness-proportional selection. That is, a number of offspring are created equal to the population size of each species where the parent of each offspring is chosen at random from the current population of that species, with probability proportional to each individual's fitness accumulated during the evaluation period. Offspring are then subjected to mutation as described previously. In this work, we use a constant mutation rate of one mutation per offspring, with all four types being equally likely. Once the new population has been generated, the previous generation is discarded.

Complexity Metric

Our goal in this work is to measure the growth of complexity of organisms over the course of an evolutionary simulation. To do so, we require a quantitative metric of complexity. A simple approach would simply be to count the number of nodes and transitions in a particular organism, and define those with more nodes to be more complex. However, consider the example of an organism with a great many nodes, all of which are labelled to produce the symbol 0. This simple metric would label this organism as very complex, but its behavior would be identical to an organism with a single node! This is clearly unsatisfactory. In order for an organism with a large network to be truly complex, it seems that it should be the case that all of its nodes are essential to its strategy.

We take inspiration from the concept of Kolmogorov complexity, which (informally) defines the complexity of a string as the length of the shortest computer program which generates that string. We define the complexity of an organism as the size of the smallest organism which produces an identical strategy. Thus, the large organism in the example above would have a very small complexity, as its strategy can be produced by an organism with a single node. Similarly, unreachable nodes will not contribute to the complexity of an organism, as those nodes could be removed without altering the strategy.

To calculate the minimal organism which produces the same strategy, we turn to automata theory. A slight modification of Hopcroft's algorithm for minimization of deterministic finite state machines (Hopcroft, 1971) allows us to efficiently compute the smallest equivalent organism. As in the base algorithm, we begin by removing all disconnected nodes. We then apply the partition-refinement process, but instead of starting with an initial partition of accepting and

non-accepting states as in Hopcroft's algorithm, we begin with a partition of nodes by output label. From there, a straightforward application of the standard algorithm produces a minimal equivalent machine. The total number of nodes in this minimized machine is then used as our complexity metric for that organism.

Experiments & Results

For each of the five ecosystems described above, we perform 20 experiments of 10,000 generations each, with each species having a population size of 50. For each generation, the complexity (the size of the minimal equivalent machine) of each organism is computed, and we record the median complexity for each population. We use the median to avoid giving undue influence to a single complex mutant which does not manage to spread its complexity throughout the population. We then take the mean of these medians over the 20 experiments to determine the characteristic trajectory of complexity over time for each ecosystem. To analyze the trajectory of complexity growth within each species at the end of the simulation, we compute a line of best fit over the complexity values for the last 2,500 generations of each experiment using the least squares method, and report the slope of this line. Species which sustained complexity growth through the end of the experiment display a positive slope, while those whose complexity growth had slowed show a near-zero slope. We will present results based on this trajectory, as well as the final level of complexity at the end of 10,000 generations. Figure 5 shows the trajectories of average complexity for each ecosystem. Figure 5 summarizes final and maximal average complexity values as well as complexity trendline slopes for each ecosystem.

Single-Species (Control)

In this ecosystem, we have only a single species which experiences no selective pressure - all organisms are equally likely to reproduce. Thus, the only driver of increased complexity is the random walk of genetic drift. Because this walk is bounded below at 1, we should expect a slight upward trend over time. This is indeed what we observe, with the populations achieving an average median complexity of about 12 nodes after 10,000 generations, with a high-water mark of 14.24 reached after 7,589 generations. By the end of the simulation, the control species had reached a near-zero rate of complexity growth. This will serve as a baseline for comparison to other systems - complexity growth above this rate should be the result of selective pressure.

Two-Species Cooperative

This ecosystem shows the behavior of two species with a purely cooperative interaction. Individuals which are able to consistently match the output of the other species are favored for reproduction. We find that this system is more amenable to complexity growth than the control - both in terms of its

final average complexity as well as the rate of growth. The two species reach final complexities of about 18 nodes after 10,000 generations, with highs of 19.78 and 22.52 reached after 6,388 and 8,507 generations respectively. By the end of the simulation, both species showed sustained, although slow, complexity growth at a rate of around 1.5 additional nodes per 1,000 generations. In some sense, this growth of complexity is surprising - the fully cooperative game can be played perfectly by two organisms with a single state, so it seems that extra states would be superfluous. From manual inspection of organism controllers, we observe that there is pressure towards strategies which are tolerant against mutations of their opponents. For example, consider a population of single-node organisms which all output 0. In each generation, there will likely be some offspring that have undergone a label-flipping mutation and instead output 1. Interaction with these mutants would result in zero fitness for the all-0 players. However, a two-state organism which always copies the last move of its opponent would achieve maximum fitness against both the original all-0 players and the all-1 mutants. Similar interactions drive further complexity development.

Two-Species Competitive

This ecosystem is the mirror of the cooperative environment above - the two species compete to achieve opposing goals: one receives fitness for matching output symbols, and the other for mismatching. Here we observe a very fast initial growth of complexity, with species reaching an average complexity of 10 nodes after only about 400 generations. By comparison, the purely cooperative system took 1,400 generations to reach this level, and the control took over 3,200 generations. However, this fast initial growth is not sustained. The system reaches final complexity values of about 16 nodes, with highs of 19.88 and 20.7, all slightly lower than those achieved by the cooperative model. The final complexity trendlines show a modestly increasing trajectory for one species (0.5 additional nodes per 1,000 generations) and a modestly decreasing trajectory for the other (-0.8 nodes per 1,000 generations). This system appears to fall victim to a form of mediocre stable state as described by (Ficici and Pollack, 1998), in which the system ceases to make progress as measured by an external metric (in our case, complexity). Instead, the species are driven into a cyclical pattern of what Ficici terms "convention-chasing". That is, species spend all of their available mutational change trying to keep up with the changes of their opponents, rather than become more complex. In particular, selection seems to favor cycles of label-changing mutations, which pushes the populations through variants of the same core strategies but with different output symbols.




Figure 5: Mean of medians complexity values for the six ecosystems over 10,000 generations. Vertical axes denote complexity as measured in nodes in the minimal equivalent strategy, and horizontal axes denote generations.

Three-Species Mixed

This ecosystem contains three species with both cooperative and competitive interactions. Unlike the previously discussed systems, the roles of the species are not symmetric. Species *A* faces both cooperative and competitive interactions, species *B* faces only a competitive interaction - its interaction is locally equivalent to those in the two-species competitive model, and species *C* faces only a cooperative interaction - locally equivalent to those in the two-species cooperative model. We observe significantly higher complexity values in this system than found in any of the two-species models. Species *A* reaches a final complexity of about 26 nodes, with a high of 29.12. Species *B* reaches a final complexity of about 20 nodes, with a high of 23.14, and species *C* reaches a final complexity of about 41 nodes, with a high of 41.72. On the other hand, the final complexity trajectories show that growth has slowed to less than one additional node per 1,000 generations for all three species by the end of the simulation. Species *C* in particular appears to suffer a period of regression around the 8,000th generation that arrests its otherwise consistent complexity growth. The

exact cause of this regression is unclear.

The complexity values for *B* and *C* are both higher than those found in the two-species competitive and cooperative games respectively, although the values for *B* only slightly so. The bulk of the increase in complexity seems to occur from the interaction between *A* and *C*, with both significantly exceeding the complexity levels observed in the purely cooperative game. The mechanism driving this growth seems to be similar to that in the cooperative game - *C* is driven to develop tolerance against changes in the behavior of *A*. However, the changes in the behavior of *A* are no longer merely the result of genetic drift, they're driven by selection pressure exerted by *B*.

Three-Species Comp

It might fairly be asked whether the increase in complexity shown in the previous system was merely due to the introduction of a third species - perhaps more players results in more complexity. To test this hypothesis, we consider a system with the same interactions, except that the game between *A* and *C* has changed to COMP. Here we find no in-

Species	Final Complexity	Max Complexity	Complexity Trend
Control A	12.44	14.24	0.006
Coop A	18.28	19.78	1.380
Coop B	18.74	22.52	1.645
Comp A	16.00	19.88	0.483
Comp B	16.02	20.70	-0.865
3-Mix A	26.18	29.12	-0.082
3-Mix B	20.54	23.14	0.718
3-Mix C	41.72	41.72	0.568
3-Comp A	13.76	16.10	0.141
3-Comp B	14.02	16.48	-0.127
3-Comp C	13.08	16.16	-0.259
4-Comp A	26.68	29.92	0.140
4-Comp B	23.82	26.22	1.374
4-Comp C	38.40	39.78	3.089
4-Comp D	36.10	37.98	3.767

Figure 6: Final and maximal complexity values, and final complexity growth slopes for each species within each ecosystem. The largest values for each column are in bold. Complexity growth values are the rate of growth in complexity per 1,000 generations, and are calculated over the final 2,500 generations.

crease in complexity above that observed in the two-player case. In fact, the convergence to mediocre stability is even stronger, with all three species quickly reaching a point of stagnant complexity. The three species reach final complexities of 13.7, 14.0 and 13.0 respectively. They reach highs of 16.1, 16.5 and 16.1, respectively. All of these values are slightly below those observed in the two-species competitive case, and well below the levels observed in the three-species mixed system. Further, all three species display final trajectory trendlines with slopes very close to zero. We therefore conclude that the interaction of competition and cooperation is driving the increase in complexity, not merely the presence of a third species.

Four-Species Mixed

The four-species mixed ecosystem is formed by adding an additional species to the three-species mixed system, such that both *A* and *B* have a cooperative partner. Here we find that *B* attains similar complexity to that of *A* in the three-species mixed system, and *D* attains similar complexity to that of *C*. In particular, *A* and *B* attain final complexities of 26.7 and 23.8 nodes, respectively, with highs of 29.9 and 26.2. *C* and *D* attain final complexities of 38.4 and 36.1 and highs of 39.8 and 38.0. We also see that species *C* and *D* sustain final complexity trajectories well above that observed in any other simulations, both with rates of more than

3 additional nodes per 1,000 generations.

Discussion

Our experiments reveal that co-evolution is able to drive complexity growth above that expected by genetic drift - indeed, all of our coevolutionary systems achieve higher final and maximal average complexity values than the control. However, the trajectories of complexity growth vary significantly depending on the nature of the coevolutionary interaction. Competitive co-evolution is able to drive extremely fast initial complexity growth, but then settles into stagnation. Cooperative co-evolution, somewhat surprisingly, is also able to drive complexity growth, but does so at a slower, steadier pace than competition.

Extension to three species systems reveals that a mix of competitive and cooperative relationships is a much more powerful driver of complexity growth than either competition or cooperation alone. Such a system seems to alleviate the phenomenon of mediocre stable states observed in two- and three-species purely competitive systems. Returning to the competing hypothesis of Gould and Dawkins, we find that genetic drift alone is able to produce complexity growth. However, this baseline growth is well below that observed in coevolutionary systems. The competitive arms races hypothesized by Dawkins are successful in creating complexity above that created by genetic drift, but fall short of other more successful systems.

Future Work

We've presented experimental results for a limited selection of the possible multi-species ecosystems. It may be worthwhile to perform a systematic analysis of all possible ecosystems of a given size, in the style of Wolfram's analysis of elementary cellular automata. The dynamics of competitive systems have recently gained importance in the field of deep learning due to the success of generative adversarial networks (GANs) (Goodfellow et al., 2014). These systems often display the undesirable behavior observed in many purely competitive coevolutionary simulations, including those presented here. Much work has been done to coax these models towards more stable dynamics (Salimans et al., 2016; Arjovsky et al., 2017). We look towards applying the notion of mixed cooperative and competitive systems to the GAN architecture. A few works have successfully implemented systems which are approximately isomorphic to our three-species mixed model (Chen et al., 2016; Abadi and Andersen, 2016), although with the aim of extending the capabilities of the system, rather than improving learning trajectories.

Acknowledgements

The authors would like to acknowledge support for this work provided by the Alfred Schonwaller graduate fellowship in 2016.

References

- Abadi, M. and Andersen, D. G. (2016). Learning to protect communications with adversarial neural cryptography. *arXiv preprint arXiv:1610.06918*.
- Adami, C. (1998). *Introduction to artificial life*, volume 1. Springer Science & Business Media.
- Adami, C., Ofria, C., and Collier, T. C. (2000). Evolution of biological complexity. *Proceedings of the National Academy of Sciences*, 97(9):4463–4468.
- Angeline, P. J., Saunders, G. M., and Pollack, J. B. (1994). An evolutionary algorithm that constructs recurrent neural networks. *IEEE transactions on Neural Networks*, 5(1):54–65.
- Arjovsky, M., Chintala, S., and Bottou, L. (2017). Wasserstein gan. *arXiv preprint arXiv:1701.07875*.
- Bedau, M. A., McCaskill, J. S., Packard, N. H., Rasmussen, S., Adami, C., Green, D. G., Ikegami, T., Kaneko, K., and Ray, T. S. (2000). Open problems in artificial life. *Artificial Life*, 6(4):363–376.
- Chen, X., Duan, Y., Houthoofd, R., Schulman, J., Sutskever, I., and Abbeel, P. (2016). Infogan: Interpretable representation learning by information maximizing generative adversarial nets. In *Advances in Neural Information Processing Systems*, pages 2172–2180.
- Dawkins, R. (2010). Universal darwinism. *The Nature of Life: Classical and Contemporary Perspectives from Philosophy and Science*, page 360.
- Dawkins, R. and Krebs, J. R. (1979). Arms races between and within species. *Proceedings of the Royal Society of London B: Biological Sciences*, 205(1161):489–511.
- Dorin, A. and Korb, K. B. (2010). Network measures of ecosystem complexity. In *ALIFE*, pages 323–328.
- Ficici, S. G. and Pollack, J. B. (1998). Challenges in coevolutionary learning: Arms-race dynamics, open-endedness, and mediocre stable states. In *Proceedings of the sixth international conference on Artificial life*, pages 238–247. MIT Press.
- Ficici, S. G., Pollack, J. B., et al. (1998). Coevolving communicative behavior in a linear pursuer-evader game. In *Proceedings of the Fifth International Conference of the Society for Adaptive Behavior*, page.
- Gibbons, R. (1992). *Game theory for applied economists*. Princeton University Press.
- Goodfellow, I., Pouget-Abadie, J., Mirza, M., Xu, B., Warde-Farley, D., Ozair, S., Courville, A., and Bengio, Y. (2014). Generative adversarial nets. In *Advances in neural information processing systems*, pages 2672–2680.
- Gould, S. J. (2011). *Full house the spread of excellence from Plato to Darwin*. Belknap Press of Harvard University Press.
- Hopcroft, J. E. (1971). An $n \log n$ algorithm for minimizing states in a finite automaton. Technical report, Stanford, CA, USA.
- Kolmogorov, A. N. (1968). Three approaches to the quantitative definition of information. *International journal of computer mathematics*, 2(1-4):157–168.
- Lindgren, K. (1992). Evolutionary phenomena in simple dynamics. In *Artificial life II*, pages 295–312.
- Mookherjee, D. and Sopher, B. (1994). Learning behavior in an experimental matching pennies game. *Games and Economic Behavior*, 7(1):62–91.
- Rissanen, J. (1978). Modeling by shortest data description. *Automatica*, 14(5):465–471.
- Salimans, T., Goodfellow, I., Zaremba, W., Cheung, V., Radford, A., and Chen, X. (2016). Improved techniques for training gans. In *Advances in Neural Information Processing Systems*, pages 2226–2234.
- Taylor, T., Bedau, M., Channon, A., Ackley, D., Banzhaf, W., Beslon, G., Dolson, E., Froese, T., Hickinbotham, S., Ikegami, T., et al. (2016). Open-ended evolution: perspectives from the oee workshop in york. *Artificial Life*.
- Wolfram, S. et al. (1984). Cellular automata as models of complexity. *Nature*, 311(5985):419–424.
- Yu, S. (2001). State complexity of regular languages. *Journal of Automata, Languages and Combinatorics*, 6(2):221.